

Urmit Mahida

+91 7795091699 | urmitmahida34@gmail.com | LinkedIn | Portfolio | Bengaluru, India

Summary

Data Scientist with 2.5 years of experience building production-grade GenAI and ML systems. Specialized in RAG pipelines, LLMs, NLP, and predictive modeling with deployments on Azure and GCP delivering measurable business impact. Research background with a published IEEE paper on AI driven optimization.

Experience

NTT Data Inc.

Jul 2024 – Mar 2026

System Integration Senior Specialist – Data Scientist, GenAI & ML

Bengaluru, India

- Engineered and productionized a **RAG pipeline** using FAISS vector indexing, Sentence Transformers (all-MiniLM-L6-v2), and Azure OpenAI GPT-4o. Deployed on Azure Databricks with a FastAPI inference layer, serving real-time triage recommendations at **80% acceptance rate**.
- Built a **Transformation Analytics pipeline** leveraging GPT-4o for structured insight extraction from unstructured ITSM data. Applied Topic Modeling and **K-Means Clustering**, boosting **label accuracy by 30%** across enterprise accounts.
- Developed an **ITSM ticket volume forecasting model** combining XGBoost and Random Forest with lag, rolling, and seasonal features; achieved **9% MAPE** and **4% improvement** over ARIMA baseline, tracked via **MLflow**.

Zee Entertainment Enterprises Ltd.

Jul 2023 – Mar 2024

Associate Data Scientist

Bengaluru, India

- Conducted RCA for KPIs including Daily Active Users, Registration, and Login success rates; drove feature optimizations improving user engagement metrics.
- Built a **BERT-based model** for user review classification (**7% accuracy improvement**) and a Text-to-SQL LLM pipeline using Gemini Pro for natural language querying of OTT databases.
- Automated acquisition dashboards and KPI reports on GCP using PySpark, enabling faster data-driven decisions across product teams.

Projects

Finsight AI Agent | Live Demo

- Engineered a **6-agent** CrewAI pipeline that fetches live NSE/BSE market data, analyses news sentiment, scores fundamentals, and generates BUY/HOLD/SELL research reports with a 12-month target price in under **90 seconds**.
- Reduced LLM calls from **40+ to 6** per analysis by pre-fetching all data in pure Python and injecting it into agent context, with same-day JSON caching for instant repeated queries. Deployed via Docker, GitHub Actions CI/CD, and SSE streaming for real-time progress updates.

Education

Indian Statistical Institute, Kolkata

Aug 2021 – Jun 2023

M.Tech in Computer Science

Ramrao Adik Institute of Technology, Navi Mumbai

Jul 2015 – Jun 2019

B.E. in Information Technology

Publications

Quantum-Inspired Constraint Optimization with Dynamic Explainable AI for Resilient EV Charging Infrastructure in Smart Cities — Accepted at IEEE Global Humanitarian Technology Conference (BHTC), 2026.

Technical Skills

Programming Languages: Python, SQL, C

Generative AI & LLMs: LLMs, RAG, LangChain, Azure OpenAI (GPT-4o), Gemini Pro, Ollama, Prompt Engineering

ML & Deep Learning: Scikit-Learn, XGBoost, Random Forest, LSTM, Transformers, BERT, TensorFlow, PyTorch, CLIP, Sentence Transformers, FAISS, NLTK, OpenCV

MLOps & Data Engineering: MLflow, PySpark, Pandas, NumPy, FastAPI, Streamlit, Matplotlib, Seaborn

Cloud & Tools: Azure Databricks, GCP, Vertex AI, BigQuery, Looker, MySQL, VS Code